

Bid Pilot

Should we bid, and if so, how do we win the score?

SMALL B2G TEAMS

The bid is lost before writing.

FRAGMENTED PURSUIT WORK

DECISION → SCORE

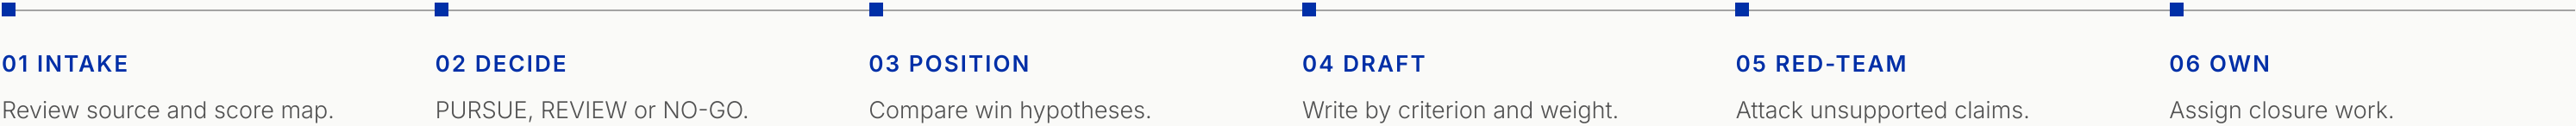
Teams pursue opportunities they cannot credibly win, then draft generic prose without connecting the buyer's scoring model to supplier evidence.

Qualification, strategy, writing, review and ownership live in separate files. A general LLM can summarize a notice. It cannot preserve an accountable decision chain.

BIDPILOT MAKES THAT CHAIN THE PRODUCT.

PRODUCT WORKFLOW

From tender to owned work.



Governed memory, not a chat transcript.

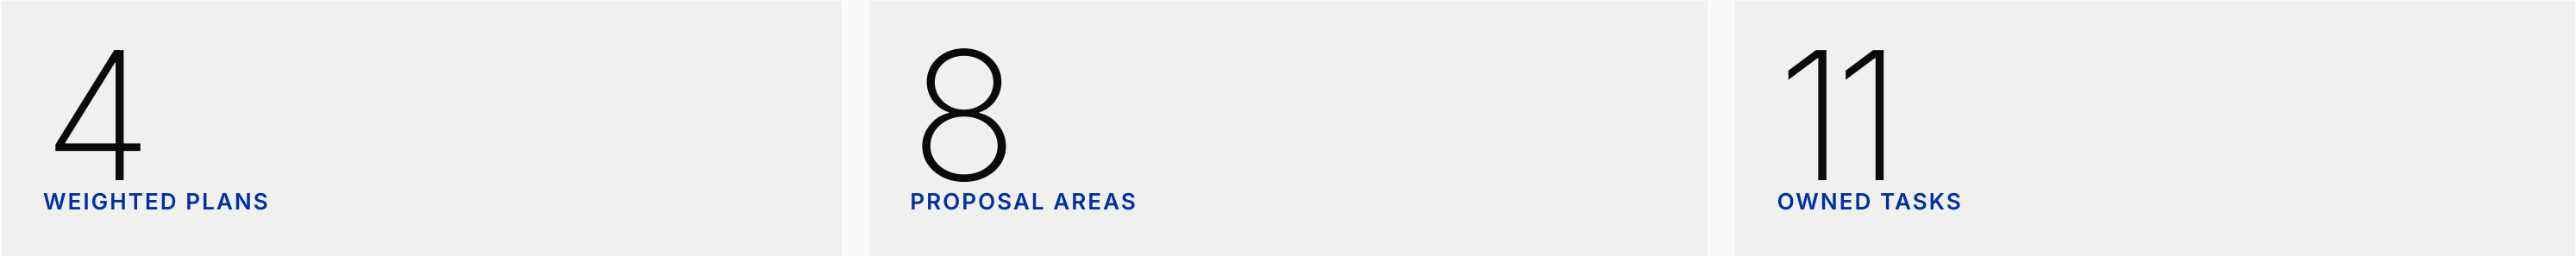
The Opportunity Graph joins versioned tender requirements with credentials, capacity, people and comparable delivery. Every artifact stays under one run ID.

CORTEX CODE CLI	Creates strategy, weighted plans, proposal areas and closure tasks.
SNOWPARK POLICY	Runs eligibility, capacity and delivery decisions next to governed data.
STREAMLIT BID ROOM	Reloads complete least-privilege runs for review and download.

VERIFIED RUN

Built inside Snowflake.

The same run is replayed through BIDPILOT_READER with no fixture fallback.



POLICY	Python ↔ Snowpark	2 × 2 parity
AGENT	Cortex Code CLI	session + query IDs
REPLAY	Authenticated reader	same run_id

- ☒ Authenticated Snowflake Bid Room
- ☐ Bid Room replay
- ☐ Tender intake
- ☐ Synthetic decision simulation

Persisted run

cortex-final-20260802-a ▼

02 Official weighted evaluation score map

100 of 100 weighted points carry a recorded supplier asset

■ Evidence recorded · 100 pts ■ No asset recorded · 0 pts

Evidence is derived in this view from RUBRIC_RESPONSE_PLANS.assets: a criterion counts as evidenced when the run recorded at least one supplier asset against it.

CRITERION	WEIGHT	EVIDENCE	SUPPLIER ASSET	PLANNED CLAIM
Technical approach Owner · Technical Lead	40	Asset recorded	project-open-data credential:Information-system maintenance certificate	Northstar applies a structured data-quality remediation methodology proven in the City Open Data Reliability Program, combining automated defect detection with API change-control gates to prevent regression.
Comparable delivery Owner · Delivery Manager	30	Asset recorded	project-open-data project-api	The City Open Data Reliability Program and Citizen API Service Transition together demonstrate delivery of data-quality and API-continuity outcomes at comparable scale in public-sector contexts.
Delivery team Owner · Resource Manager	20	Asset recorded	availability:900h project-open-data project-api	Northstar provides a dedicated team with 900 available hours, exceeding the 720-hour requirement, drawn from staff who executed the City Open Data and Citizen API programs.
Price Owner · Commercial Lead	10	Asset recorded	availability:900h	Pricing is anchored to the 720-hour delivery envelope with no subcontracting margin, enabled by the capacity surplus that eliminates overtime premiums.

JUDGE-FACING PRODUCT

The 40-point moment.

DECISION

PURSUE is the permission boundary.

SCORE MAP

Official weight controls priority.

EVIDENCE

Every claim names an asset and owner.

VERIFIED REVERSALS

Fluent fiction stays locked.

P

PURSUE

All gates pass.

proposal open

R

REVIEW

Evidence needs validation.

proposal locked

N

NO-GO

Eligibility or capacity fails.

proposal locked

O

INVENTED FACTS

Missing people, pricing and metrics become tasks.

never asserted

Win the score, not the prompt

01

Decide

Stop wasted pursuits before drafting.

02

Position

Turn official weights and proof into strategy.